

R3 Waiting in Line

14:540:384: Simulation Models in IE (Spring 2025)

Daniel Moore

2025-02-04

This document shows some basics about how to leverage Quarto to make your life easier and how to create some neat animations.

Table of contents

Recap	2
Learning Objectives	2
1 Quarto Controls	2
1.1 Front Matter	2
1.2 Cell Options	3
1.3 Formulas and references	4
2 Bank Queuing Simulation	4
2.1 Problem Definition	4
2.2 Coding the simulation	5
2.2.1 Packages	5
2.2.2 Parameters	5
2.2.3 Generating Arrival and Service Times	6
2.2.4 Running the Simulation	6
2.2.5 Calculating Performance Measures	6
2.3 Visualizing the Results	8
2.3.1 Queue Length Histogram	8
2.3.2 Waiting Time Histogram	9
2.3.3 Waiting Time Density	10
2.3.4 Queue Length Over Time	11
3 Animation	15
3.1 Concepts	15

3.2	Code	15
4	Analytical Solution	16
4.1	Formulae	16
4.2	Calculations	17
5	Simulation and Analytical Comparison	17
5.1	Think Critically!	18
5.2	Performance Measure Notes	18
5.3	What to do?	18
6	Conclusion	20
6.1	Potential Complications?	20
6.2	Summary	20

Questions

Notice this was excluded from the Table of Contents and is not numbered.

Recap

- Observed how simple problems can require simulation
- Set up a Quarto document for R, downloaded data, and generated some visualizations
- Fit various distributions to data and evaluated the best fit

Learning Objectives

- Control output and content of Quarto cells
- Create and run a basic queueing simulation
- Plot results
- Create an animation of the queue process

1 Quarto Controls

1.1 Front Matter

```

---
title: "R3 Waiting in Line"
subtitle: "14:540:384: Simulation Models in IE (Spring 2025)"
author:
  - name: Daniel Moore
    email: daniel.l.moore@rutgers.edu
    affiliation:
      - name: Rutgers University
        city: Piscataway
        state: NJ
        url: https://sites.rutgers.edu/ropes-lab/

date: 2025-02-04
image: "../assets/queue_plot.jpeg"
date-format: iso

format:
  pdf:
    number-sections: true
    toc: true
    toc-depth: 3
    number-depth: 3
    output-file: R3_pdf
    df-print: kable

abstract: This document shows some basics about how to leverage Quarto to make your life eas
---

```

1.2 Cell Options

```

#| eval: true      # whether to execute code when you render the doc
#| echo: false     # decide if code is shown in the document or hidden
#| output: false   # should the output of the chunk be displayed
#| warning: false  # do you want warnings to be shown in the document
#| error: false    # do you want error messages to appear in the doc?
#| include: false  # Catch all for preventing any output

```

1.3 Formulas and references

- create a cross referencable item with `#fig-your-reference-name`. Use the appropriate *reserved* prefix such as: `fig`, `tbl`, `lst`, `tip`, `nte`, `wrn`, `imp`, `cau`, `thm`, `lem`, `cor`, `prp`, `cnj`, `def`, `exm`, `exr`, `sol`, `rem`, `eq`, `sec`.
- link words in your reference with `-`. Do not using blank space, `,` or underscores `_`.
- Reference your item where you want with `@fig-your-reference-name`.
- See the [Quarto References Guide](#) for how to get different formats of your reference such as: “Figure 1”, “Fig 1”, “1” etc.

The probability density function (pdf) of the exponential distribution is given below in `@eq-exponential-distribution`.

```
$$  
p(t) = \lambda e^{-\lambda t}  
$$ {#eq-exponential-distribution}
```

The probability density function (pdf) of the exponential distribution is given below in Equation 1.

$$p(t) = \lambda e^{-\lambda t} \tag{1}$$

2 Bank Queuing Simulation

2.1 Problem Definition

From 4:00 PM to 7:00 PM, customers arrive at Whole Foods to make Amazon returns following a Poisson process with a rate of $\lambda = \frac{1 \text{ customer}}{3 \text{ min}}$. A single clerk services the returns at a rate of $\mu = \frac{1 \text{ customer}}{2 \text{ min}}$.

What can we say about the:

- Expected wait time?
- Expected total time in the system?
- What is the probability of seeing more than 5 people in line?
- Server Utilization Rate?

2.2 Coding the simulation

- For previous examples, we stepped through time and sampled the random variables, keeping track of the results
- For this example, we will sample all the random variables and perform some logic to determine the results
- This method is more computationally efficient because loops can slow down R and Python
- On the one hand, sampling the random variables this way is more convincing because the memoryless property of the exponential distribution can be a hard concept to grasp and we are kind of pushing that issue out of the way
- On the other hand, this method requires a little more reasoning about what is happening in the system after the random variables are sampled

2.2.1 Packages

```
library(tidyverse)
library(knitr)

# setting a random seed for reproducibility
set.seed(1766)
```

2.2.2 Parameters

```
lambda <- 1/3
arrive <- function(n=1) {
  rexp(n, rate = lambda )
}

mu <- 1/2
service <- function(n=1) {
  rexp(n, rate = mu)
}

t_f <- 3 * 60

# double the number of samples to make sure we will have enough to exceed t_f
samples <- round(2 * t_f * lambda)
```

2.2.3 Generating Arrival and Service Times

```
data <- tibble(  
  interarrival_time = arrive(samples),  
  service_time = service(samples)  
)  
  
data <- data |> mutate(  
  arrival_time = cumsum(interarrival_time)  
)  
  
data <- data |> filter(arrival_time <= t_f)  
  
data <- data |> mutate(  
  service_start_time = 0,  
  departure_time = 0)
```

2.2.4 Running the Simulation

```
t_0 <- 0  
next_available_time <- t_0  
  
for (i in 1:nrow(data)) {  
  # Service starts when the server is available  
  data$service_start_time[i] <- max(data$arrival_time[i], next_available_time)  
  
  # Compute departure time  
  data$departure_time[i] <- data$service_start_time[i] + data$service_time[i]  
  
  # Update when the server will next be available  
  next_available_time <- data$departure_time[i]  
}
```

2.2.5 Calculating Performance Measures

```
data <- data |> mutate(  
  wait_time = service_start_time - arrival_time,  
  total_time = departure_time - arrival_time
```

```

)

data <- data |> mutate(
  number_served_by_arrival = 0,
  L = 0, # number in system, waiting + being serviced. Starting at 0
  L_q = 0 # number in queue
)

for (i in 2:nrow(data)) {
  this_arrival <- data$arrival_time[i]
  data$number_served_by_arrival[i] <- sum(data$departure_time[1:i-1] < this_arrival)
}

for (i in 1:nrow(data)) {
  data$L[i] <- i - data$number_served_by_arrival[i]
  data$L_q[i] <- max(0, data$L[i] - 1)
}

head(data)

```

interarrival_time	arrival_time	service_time	service_start_time	departure_time	wait_time	total_time	number_served	L	L_q	arrival
0.6624587	5.2966049	0.6624587	0.6624587	5.959064	0.0000000	6.624587	0	1	0	
9.3129468	2.5289920	0.9754055	5.9754055	12.504397	0.0000000	10.528992	1	1	0	
0.3265047	2.9688180	0.3019102	6.2504397	15.473216	2.2024873	11.713053	1	2	1	
0.2567624	2.1481199	0.5586726	6.7547321	17.621335	4.9145429	10.626628	1	3	2	
0.5007908	1.0292805	1.0594634	7.7621335	18.650616	6.5618779	9.5911524	1	4	3	
7.1882326	0.4291179	1.8247696	8.6506159	19.079734	0.4029108	8.320377	4	2	1	

Note that this table will look bad in the PDF. We'll explicitly use `kable()` to make it look better. Look at Table 2 to see compare the outputs.

```

kable(head(data), digits=2, col.names =
  c("Interarrival Time", "Arrival Time", "Service Time", "Service Start Time", "Departure Time", "Wait Time", "Total Time", "Number Served", "L", "L_q", "Arrival Time"))

```

Table 2: Data Formatted with Kable

Interarrival Time	Arrival Time	Service Time	Service		Wait Time	Total Time	Serviced by Arrival	L	L_q
			Start Time	Departure Time					
0.66	5.30	0.66	0.66	5.96	0.00	5.30	0	1	0
9.31	2.53	9.98	9.98	12.50	0.00	2.53	1	1	0
0.33	2.97	10.30	12.50	15.47	2.20	5.17	1	2	1
0.26	2.15	10.56	15.47	17.62	4.91	7.06	1	3	2
0.50	1.03	11.06	17.62	18.65	6.56	7.59	1	4	3
7.19	0.43	18.25	18.65	19.08	0.40	0.83	4	2	1

We could also break this table up into a few different groups of columns so that it fits on the page nicer. This is good enough for us for now.

2.3 Visualizing the Results

2.3.1 Queue Length Histogram

Figure 1 shows the distribution of queue lengths.

```
data |> ggplot(aes(x=L)) +
  geom_histogram(fill="cornflowerblue", color="black") +
  labs(
    title = "Queue Length Histogram",
    x = "Number of Customers in Queue",
    y = "Count"
  )
```

`stat_bin()` using `bins = 30`. Pick better value with `binwidth`.

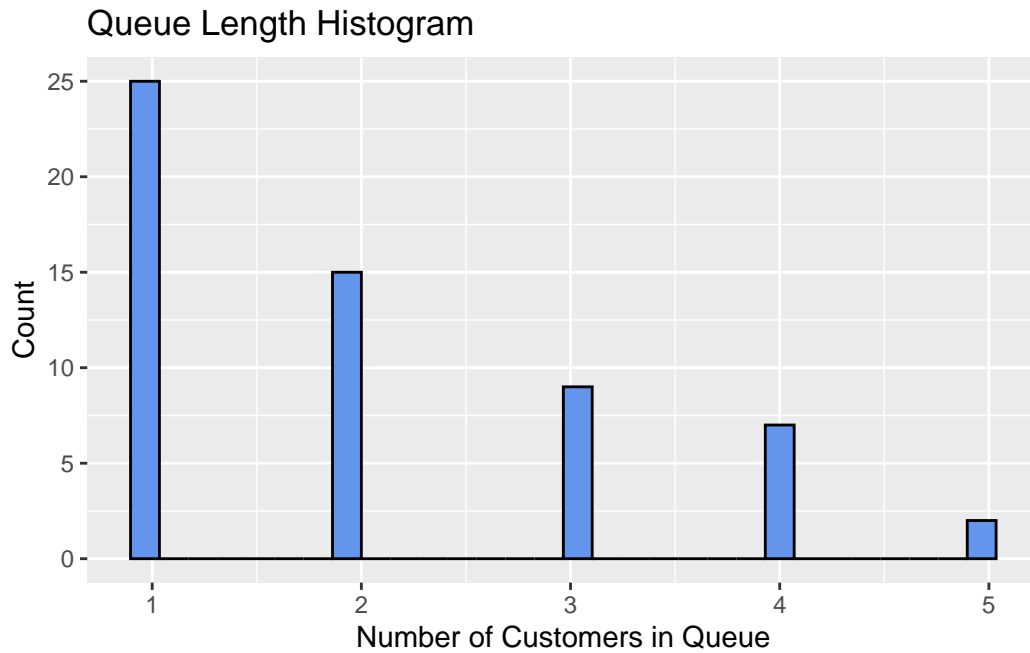


Figure 1: Histogram of Queue Lengths

2.3.2 Waiting Time Histogram

Figure 2 shows the distribution of waiting times.

```
data |> ggplot(aes(x=wait_time)) +  
  geom_histogram(fill="cornflowerblue", color="black") +  
  labs(  
    title = "Waiting Time Histogram",  
    x = "Waiting Time (min)",  
    y = "Count"  
  )
```

``stat_bin()`` using ``bins = 30``. Pick better value with ``binwidth``.




Figure 2: Histogram of Waiting Times

2.3.3 Waiting Time Density

If we want to view any of these as a probability density we could use `stat_density` or `geom_density`. There are a whole [host of different plots that you should browse](#) to see what could be good at telling the story you are trying to convey.

```
data |> ggplot(aes(x = wait_time)) +  
  geom_density() +  
  labs(  
    x = "Time in the System (min)",  
    y = "Probability Density"  
  )
```

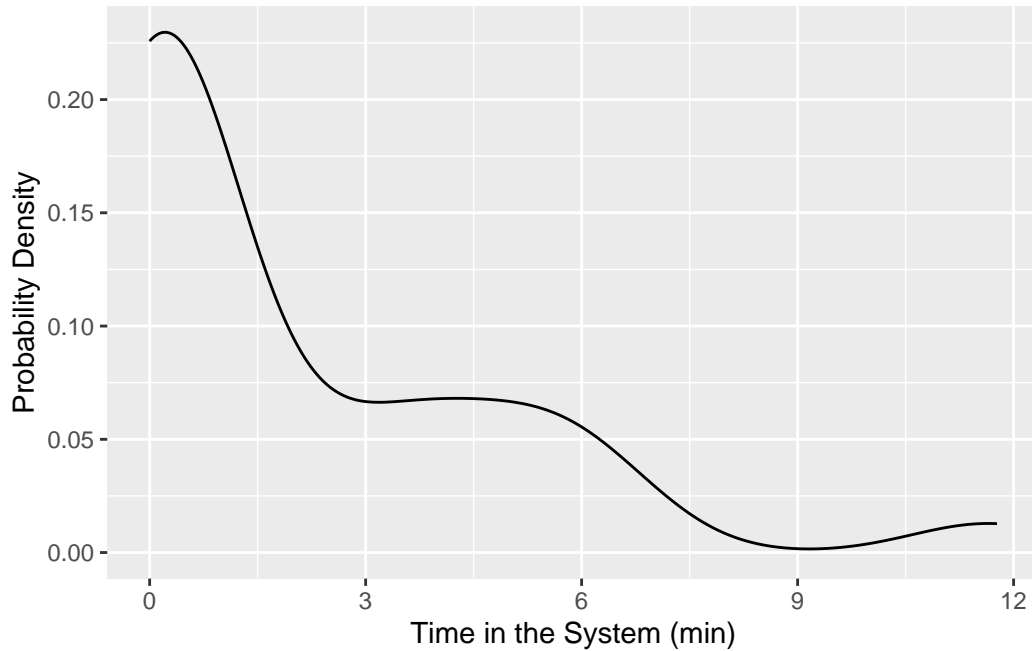


Figure 3: System Waiting Time Density Plot

2.3.4 Queue Length Over Time

Figure 5 shows the queue length over time. However, this isn't really the full story. What we want to see is how the system changes at each arrival and departure. To do this, we will need to create a new data frame that has an entry for each arrival and departure. Additionally, we'll want the numbers indicated to be horizontal lines with discontinuities showing when the number changes. Also, this is not exactly the bounds of our simulation because we cutoff our arrivals at t_f , but we have departures after that time. So we should only be plotting up to t_f .

```
data |> ggplot(aes(x=arrival_time, y=L_q)) +
  geom_line(color="cornflowerblue") +
  labs(
    title = "Queue Length Over Time",
    x = "Time (min)",
    y = "Number of Customers in Queue"
  )
```




Figure 4: Queue Length Over Time

The way these lines are connected give an indication that the queue length is in between two integers at most times. Line plots can be misleading in this way. We need a step plot. Also, we are essentially only sampling the system at a time that somebody arrives. This is why we are seeing big jumps. We should only be seeing the line move up or down one at a time.

Event Log

Basically, we need to unpack the simulation data so that it is event-centric rather than customer-centric.

```
# all events are either arrivals or departures
# create a dataframe which tracks all events
# combines arrival and departure times into single vector, event time
# for arrivals, event is +1
# for departures, event is -1

events <- data |>
  # bringing the data I want to work with
  select(arrival_time, departure_time) |>

  # turning the data from wide to long
```

```

pivot_longer(
  cols = everything(),
  names_to = "event_type",
  values_to = "event_time") |>

# assigning -1 if departure and 1 if arrival
mutate(
  event = ifelse(event_type == "arrival_time", 1, -1)) |>

# adding an initial condition of time=0 and 0 customers
bind_rows(tibble(event_time = 0, event = 0)) |>

# removing the event type column
select(-event_type) |>

# sorting the data by event time
arrange(event_time)

# cumulative sum of events
events <- events |> mutate(
  L = cumsum(event),
  # pmax is vectorized function that looks at each element rather than the entire thing
  L_q = pmax(0, L - 1)
)

head(events)

```

event_time	event	L	L_q
0.0000000	0	0	0
0.6624587	1	1	0
5.9590636	-1	0	0
9.9754055	1	1	0
10.3019102	1	2	1
10.5586726	1	3	2

We could plot this by adding each column to the plot one at a time, but this can get messy. Instead we'll do a `pivot_longer` so that the data is in a better format for plotting multiple series together.

Here I'm pivoting from a wide format (have both an `L` and an `L_q` column) to a long format which just has the x-value of `time` and the y-value of `L` or `L_q`. This makes plotting much easier.

```
events_long <- events |>
  pivot_longer(
    cols = c(L, L_q),
    names_to = "metric",
    values_to = "value"
  )

head(events_long)
```

event_time	event	metric	value
0.0000000	0	L	0
0.0000000	0	L_q	0
0.6624587	1	L	1
0.6624587	1	L_q	0
5.9590636	-1	L	0
5.9590636	-1	L_q	0

At last, Figure 5 shows a complete plot of the number of customers in line and the number in the system at each moment.

```
events_long |> ggplot(aes(x = event_time, y = value, color=metric)) +
  xlim(0, t_f) +
  geom_step(direction = "hv") +
  labs(
    title = "Customer Count Over Time",
    x = "Time (min)",
    y = "Customers"
  ) +
  theme_minimal()
```




Figure 5: Customer Count Over Time

3 Animation

3.1 Concepts

- Data in long format: Structure data with one row per observation and a time variable for animations (use `pivot_longer()` for multiple metrics).
- Time variable for transitions: Use a continuous variable (e.g., `event_time`) to animate data over time with `transition_reveal()` or `transition_states()`.
- Compatible geom layers: Use `geom_step()`, `geom_line()`, or `geom_point()` to animate time-series data, ensuring smooth transitions.

3.2 Code

```
library(gganimate)
library(gifski)
```

Note that I have this cell set up to not execute with `#| eval: false`. I used this exact code to create and save the plot as shown, but I don't want to re-render and save over it every time.

```

p <- events_long |>
  ggplot(aes(x = event_time, y = value, color = metric)) +
  geom_step(size = 1.2) +
  labs(x = "Time", y = "Customers", color = "Metric") +
  theme_minimal() +
  transition_reveal(event_time)

# Animate and save as GIF
anim <- animate(p, renderer = gifski_renderer(),
  width = 800, height = 600, duration = 5)
anim_save("returns_animation.gif", animation = anim)

anim

```

Obviously, we can't put a gif into a pdf. Instead, it's better to insert a still image or omit it all together. See how to control how content is displayed based on the format in the [Quarto docs](#).

4 Analytical Solution

These metrics are the typical things we want to know about queues. How long will I wait? What's the probability of the line having more than 4 people? How long will I be in this line before it's my turn? etc.

They are derived from algebraic manipulations of the probability density functions. They are specific for an $M/M/1$ queue which is “[Kendall's Notation](#)” for a queue with:

- M: Exponential distributed Arrival Time (typically λ)
- M: Exponentially distributed Service Time (typically μ)
- 1: A single line, single server
- Implied is a first-in/first-out (FIFO) “service discipline”

What other service disciplines can you imagine? What settings?

4.1 Formulae

$$\text{Server Utilization, } \rho = \frac{\lambda}{\mu} = \frac{1/3}{1/2} \quad (2)$$

$$\text{Expected Number in the System, } L = \frac{\lambda}{\mu - \lambda} \quad (3)$$

$$\text{Expected Number in the Queue, } L_q = \frac{\lambda^2}{\mu(\mu - \lambda)} \quad (4)$$

$$\text{Expected time in the System, } W = \frac{1}{\mu - \lambda} \quad (5)$$

$$\text{Expected time in the Queue, } W_q = \frac{\lambda}{\mu(\mu - \lambda)} \quad (6)$$

$$\text{Probability of empty system, } P_0 = 1 - \rho \quad (7)$$

$$\text{Probability of } n \text{ customers in the System, } P(n) = (1 - \rho)\rho^n \quad (8)$$

4.2 Calculations

```
#
P_n <- function(n, rho) {
  (1 - rho) * rho^n
}

performance_measures <- data.frame(
  Performance_Measure = c("L", "L_q", "W", "W_q", "rho"),
  Analytical = c(
    lambda / (mu - lambda),
    (lambda^2) / (mu * (mu - lambda)),
    1 / (mu - lambda),
    lambda / (mu * (mu - lambda)),
    lambda / mu
  )
)
```

5 Simulation and Analytical Comparison

- How can we get these same performance measures from the simulated data?

5.1 Think Critically!

Think for a minute about the way our two data frames are structured. Is all of the information there? What do we need to consider to get it.

5.2 Performance Measure Notes

- The performance measures, L and L_q are supposed to represent the expected number of customers in the system and in the queue
- The first data is customer centric. All we see is how many people are in line when the customer shows up. W and W_q are obtainable from this.
- The second is event centric. We see the time that an event happens and the customers at that time.
- This means essentially we are just sampling at these points. It *might* be a good representation of the number in the system, but it might not be. Notably, we could definitely have systematic error because this “sampling” isn’t even random, it is always as events

5.3 What to do?

- We need to take the time-weighted average of the number of people in the system and in line
- This data is available in the `events` data frame, but it will take a little more manipulation:

```
events <- events |> mutate(  
  time_in_state = coalesce(lead(event_time), t_f) - event_time  
)
```

Coalesce gives us a default value to use. Since we are differencing each row with the one after it, when we get to the last row there’s nothing there. So we need to say the last row’s lead even time is the end of the simulation, `t_f`.

```
performance_measures <- performance_measures |> mutate(  
  Simulation = c(  
    sum(events$L * events$time_in_state) / t_f,  
    sum(events$L_q * events$time_in_state) / t_f,  
    mean(data$total_time),  
    mean(data$wait_time),  
    sum(events$time_in_state[events$L > 0]) / t_f  
  )  
)
```

Table 6: Calculating $P(n)$

```
n_in_system_ps <- tibble(n = 0:10) |>
  mutate(
    Analytical = P_n(0:10, lambda/mu),
    Simulation = sapply(0:10, function(k) sum(events$time_in_state[events$L == k]) / t_f)
  )

kable(n_in_system_ps, digits=3)
```

Now we can compare the analytical results in Table 5 and the correctly calculated simulation performance measures.

```
kable(performance_measures, digits = 3)
```

Table 5: Simulation and Analytical Comparison

Performance_Measure	Analytical	Simulation
L	2.000	1.190
L_q	1.333	0.643
W	6.000	3.692
W_q	4.000	1.995
rho	0.667	0.547

```
# first we need to get the data into long format

n_in_system_ps |>
  pivot_longer(
    cols = c(Analytical, Simulation),
    names_to = "method",
    values_to = "p"
  ) |>
  ggplot(aes(x = n, y = p, fill = method, group = method)) +
  geom_bar(stat = "identity", position = "dodge") +
  labs(x = "n", y = "Probability", title = "Probabilities by Method") +
  theme_minimal()
```

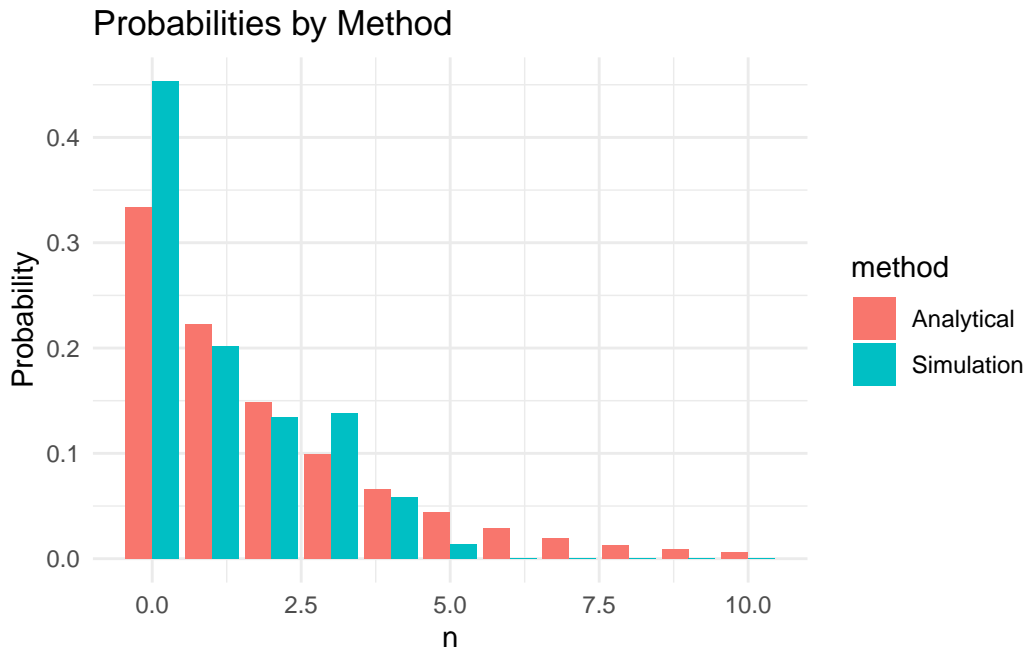


Figure 6: Corrected $P(n)$ Histogram

I'll leave it to you to compare these results with an incorrect method which would be simply averaging from the sampled data or the events log without accounting for time.

6 Conclusion

6.1 Potential Complications?

- What if there is a limited amount of space and if the line is too long, customers just leave?
- Service times follow different distributions?
- More than one line? More than one server?
- Priority line with different logic? Maybe a commercial customer with 12 packages gets serviced as soon as a server is available?

6.2 Summary

- We simulated an $M/M/1$ queuing system by sampling arrival and service times
- We obtained system state measurements by executing logic on the samples
- Manipulated the data into an event time series

- Created and animated `geom_step` plots using `ggplot`, `gganimate`, and `gifski`
- Compared analytical system metrics to those we obtained from the simulation.
- Noted how we must account for time to get correct performance measures